



CS 443 Hash-based Indexes Chapter 11

Slides adapted from
Ramakrishnan & Gerhke
pages.cs.wisc.edu/~dbbook/

Renée J. Miller

Introduction

2

- As for any index, 3 alternatives for data entries k^* :
 1. Data record with key value k
 2. $\langle k, \text{rid of data record with search key value } k \rangle$
 3. $\langle k, \text{list of rids of data records with search key } k \rangle$
- Choice orthogonal to the indexing technique
- Hash-based indexes are best for equality selections. Cannot support range searches.
- Static and dynamic hashing techniques exist; trade-offs similar to ISAM vs. B+ trees.


11/16/11



Static Hashing

3

- ❖ # primary pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- ❖ $h(k) \bmod M = \text{bucket to which data entry with key } k \text{ belongs. (} M = \text{\# of buckets)}$



Static Hashing (Contd.)

4

- ❖ Buckets contain *data entries*.
- ❖ Hash fn works on *search key* field of record r . Must distribute values over range $0 \dots M-1$.
 - e.g., $h(\text{key}) = (a * \text{key} + b) \bmod M$
 - a and b are constants; lots known about how to tune h .
- ❖ Long overflow chains can develop and degrade performance.
 - *Extendible* and *Linear Hashing*: Dynamic techniques to fix this problem.

11/16/11



Extendible Hashing

5

- ❖ Situation: Bucket (primary page) becomes full. Why not re-organize file by *doubling* # of buckets?
 - Reading and writing all pages is expensive!
 - *Idea*: Use directory of pointers to buckets, double # of buckets by *doubling the directory*, splitting just the bucket that overflowed!
 - Directory much smaller than file, so doubling it is much cheaper. Only one page of data entries is split. *No overflow page!*
 - Trick lies in how hash function is adjusted!


11/16/11



Example

6

- Directory is array of size 4.
- To find bucket for r , take last 'global depth' # bits of $h(r)$; we denote r by $h(r)$.
 - If $h(r) = 5 = \text{binary } 101$, it is in bucket pointed to by 01.




- ❖ **Insert**: If bucket is full, split it (allocate new page, re-distribute).
- ❖ If necessary, double the directory. (As we will see, splitting a bucket does not always require doubling; we can tell by comparing *global depth* with *local depth* for the split bucket.)

11/16/11



Insert $h(r)=20$ (Causes Doubling)

7



11/16/11



Points to Note

8

- ❖ $20 = \text{binary } 10100$. Last 2 bits (00) tell us r belongs in A or A2. Last 3 bits needed to tell which.
 - *Global depth of directory*: Max # of bits needed to tell which bucket an entry belongs to.
 - *Local depth of a bucket*: # of bits used to determine if an entry belongs to this bucket.
- ❖ When does bucket split cause directory doubling?
 - Before insert, *local depth* of bucket = *global depth*. Insert causes *local depth* to become $>$ *global depth*; directory is doubled by *copying it over* and 'fixing' pointer to split image page. (Use of least significant bits enables efficient doubling via copying of directory!)

11/16/11



Comments: Extendible Hashing

- ❖ If directory fits in memory, equality search answered with one disk access; else two.
 - 100MB file, 100 bytes/rec, 4K pages contains 1,000,000 records (as data entries) and 25,000 directory elements; chances are high that directory will fit in memory.
 - Directory grows in spurts, and, if the distribution of *hash values* is skewed, directory can grow large.
 - Multiple entries with same hash value cause problems!
- ❖ **Delete:** If removal of data entry makes bucket empty, can be merged with 'split image'. If each directory element points to same bucket as its split image, can halve directory.


11/16/11



Example of Linear Hashing

10

- ❖ On split, $h_{Level+1}$ is used to re-distribute entries.
- ❖ Insert 43.



Linear Hashing

11

- ❖ Directory avoided in LH by using overflow pages, and choosing bucket to split round-robin.
 - Splitting proceeds in 'rounds'. Round ends when all N_R initial (for round R) buckets are split. Buckets 0 to $Next-1$ have been split; $Next$ to N_R yet to be split.
 - Current round number is $Level$.
 - **Search:** To find bucket for data entry r , find $h_{Level}(r)$:
 - If $h_{Level}(r)$ in range 'Next to N_R ', r belongs here.
 - Else, r could belong to bucket $h_{Level}(r)$ or bucket $h_{Level}(r) + N_R$; must apply $h_{Level+1}(r)$ to find out.

11/16/11



Linear Hashing (Contd.)

12

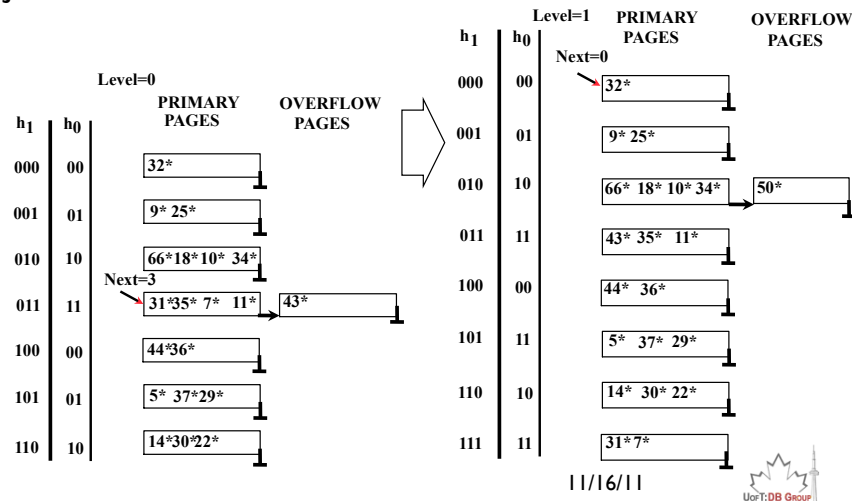
- ❖ **Insert:** Find bucket by applying $h_{Level} / h_{Level+1}$:
 - If bucket to insert into is full:
 - Add overflow page and insert data entry.
 - (Maybe) Split $Next$ bucket and increment $Next$.
- ❖ Can choose any criterion to 'trigger' split.
- ❖ Since buckets are split round-robin, long overflow chains don't develop!
- ❖ Doubling of directory in Extendible Hashing is similar; switching of hash functions is *implicit* in how the # of bits examined is increased.

11/16/11



Example: End of a Round


13



Overview of LH File

14

❖ In the middle of a round.



LH Described as a Variant of EH

15

- ❖ The two schemes are actually quite similar:
 - Begin with an EH index where directory has N elements.
 - Use overflow pages, split buckets round-robin.
 - First split is at bucket 0. (Imagine directory being doubled at this point.) But elements $\langle 1, N+1 \rangle$, $\langle 2, N+2 \rangle$, ... are the same. So, need only create directory element N , which differs from 0, now.
 - When bucket 1 splits, create directory element $N+1$, etc.
- ❖ So, directory can double gradually. Also, primary bucket pages are created in order. If they are *allocated* in sequence too (so that finding i 'th is easy), we actually don't need a directory! Voila, LH.

11/16/11

Summary

16

- ❖ Hash-based indexes: best for equality searches, cannot support range searches.
- ❖ Static Hashing can lead to long overflow chains.
- ❖ Extendible Hashing avoids overflow pages by splitting a full bucket when a new data entry is to be added to it.
 - ❖ Directory to keep track of buckets, doubles periodically.
 - Can get large with skewed data; additional I/O if this does not fit in main memory.

11/16/11

Summary (Contd.)

17

- ❖ Linear Hashing avoids directory by splitting buckets round-robin, and using overflow pages.
 - Overflow pages not likely to be long.
 - Space utilization could be lower than Extendible Hashing, since splits not concentrated on 'dense' data areas.
 - Can tune criterion for triggering splits to trade-off slightly longer chains for better space utilization.
- ❖ For hash-based indexes, a *skewed* data distribution is one in which the *hash values* of data entries are not uniformly distributed!

11/16/11

