



CS 443

Query Optimization

Chapter 15

Slides adapted from
Ramakrishnan & Gerhke
pages.cs.wisc.edu/~dbbook/

Renée J. Miller

Query Evaluation

2

□ How could we evaluate the following query?

- ▣ select Date
- ▣ from Reserves R, Sailors S
- ▣ where R.SID = S.SID and S.Rating = 10

▣ R = Reserves[Day, BID, SID]

▣ S = Sailors[SID, Name, Rating, Age]

▣ $\pi_{\text{Date}}(\sigma_{R.SID=S.SID \text{ and } Rating=10} (R \times S))$

▣ other options?

11/16/11



Some Common Techniques

3

□ Algorithms for evaluating relational operators use some simple ideas extensively:

- ▣ Indexing: Can use WHERE conditions to retrieve small set of tuples (selections, joins)
- ▣ Iteration: Sometimes, faster to scan all tuples even if there is an index. (And sometimes, we can scan the data entries in an index instead of the table itself.)
- ▣ Partitioning: By using sorting or hashing, we can partition the input tuples and replace an expensive operation by similar operations on smaller inputs.

* Watch for these techniques as we discuss query evaluation!

11/16/11



Highlights of System R Optimizer

4

□ Impact:

- ▣ Most widely used currently; works well for < 10 joins.

□ Cost estimation: Approximate art at best.

- ▣ Statistics, maintained in system catalogs, used to estimate cost of operations and result sizes.
- ▣ Considers combination of CPU and I/O costs.

□ Plan Space: Too large, must be pruned.

- ▣ Only the space of *left-deep plans* is considered.

▣ Left-deep plans allow output of each operator to be *pipelined* into the next operator without storing it in a temporary relation.

- ▣ Cartesian products avoided.

11/16/11



Overview of Query Optimization

5

- Plan: Tree of R.A. ops, with choice of alg for each op.
 - Each operator typically implemented using a 'pull' interface: when an operator is 'pulled' for the next output tuples, it 'pulls' on its inputs and computes them.
- Two main issues:
 - For a given query, what plans are considered?
 - Algorithm to search plan space for cheapest (estimated) plan.
 - How is the cost of a plan estimated?
- Ideally: Want to find best plan. Practically: Avoid worst plans!
- We will study the System R approach.



Schema for Examples

6

Sailors (sid: integer, sname: string, rating: integer, age: real)
 Reserves (sid: integer, bid: integer, day: dates, rname: string)

- Similar to old schema; rname added for variations.
- Reserves:
 - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- Sailors:
 - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

11/16/11



Query Blocks: Units of Optimization

7

- An SQL query is parsed into a collection of query blocks, and these are optimized one block at a time.
- Nested blocks are usually treated as calls to a subroutine, made once per outer tuple. (This is an over-simplification, but serves for now.)
- ❖ For each block, the plans considered are:
 - All available access methods, for each reln in FROM clause.
 - All left-deep join trees (i.e., all ways to join the relations one-at-a-time, with the inner reln in the FROM clause, considering all reln permutations and join methods.)

```
SELECT S.sname
FROM Sailors S
WHERE S.age IN
  (SELECT MAX (S2.age)
   FROM Sailors S2
   WHERE S.Rating=S2.Rating)
```

Outer block

Nested block

11/16/11



Relational Algebra Equivalences

8

- Permits the choice of different join orders and to 'push' selections and projections ahead of joins.
- Selections: $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots \sigma_{cn}(R))$
 (Cascade) $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (Commutate)
- ❖ Projections: $\pi_A(R) \equiv \pi_A(\dots(\pi_{ABC}(R)))$ (Cascade)
- ❖ Joins: $R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$ (Associative)
 $(R \bowtie S) \equiv (S \bowtie R)$ (Commutate)

☞ Show that:

$$R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$$

11/16/11



More Equivalences

9

- A projection commutes with a selection that only uses attributes retained by the projection.
- Selection between attributes of the two arguments of a cross-product converts cross-product to a join.
- A selection on just attributes of R commutes with $R \bowtie S$. (i.e., $\sigma(R \bowtie S) \equiv \sigma(R) \bowtie S$)
- Similarly, if a projection follows a join $R \bowtie S$, we can 'push' it by retaining only attributes of R (and S) that are needed for the join or are kept by the projection.

11/16/11



Enumeration of Alternative Plans

10

- There are two main cases:
 - Single-relation plans
 - Multiple-relation plans
- For queries over a single relation, queries consist of a combination of selects, projects, and aggregate ops:
 - Each available access path (file scan / index) is considered, and the one with the least estimated cost is chosen.
 - The different operations are essentially carried out together (e.g., if an index is used for a selection, projection is done for each retrieved tuple, and the resulting tuples are *pipelined* into the aggregate computation).

11/16/11



Statistics and Catalogs

11

- Need information about the relations and indexes involved. **Catalogs** typically contain at least:
 - # tuples (NTuples) and # pages (NPages) for each relation.
 - # distinct key values (NKeys) and NPages for each index.
 - Index height, low/high key values (Low/High) for each tree index.
- Catalogs updated periodically.
 - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- More detailed information (e.g., histograms of the values in some field) are sometimes stored.

11/16/11



Cost Estimation

12

- For each plan considered, must estimate cost:
 - Must estimate *cost* of each operation in plan tree.
 - Depends on input cardinalities.
 - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
 - Must also estimate *size of result* for each operation in tree!
 - Use information about the input relations.
 - For selections and joins, assume independence of predicates.

11/16/11



Cost Estimates for Single-Relation Plans

13

- Index I on primary key matches selection:
 - Cost is $Height(I) + I$ for a B+ tree, about 1.2 for hash index.
- Clustered index I matching one or more selects:
 - $(NPages(I) + NPages(R))$ * product of RF's of matching selects.
- Non-clustered index I matching one or more selects:
 - $(NPages(I) + NTuples(R))$ * product of RF's of matching selects.
- Sequential scan of file:
 - $NPages(R)$.

☞ **Note:** Typically, no duplicate elimination on projections!
(Exception: Done on answers if user says DISTINCT.)



Access Paths

14

- ❖ An access path is a method of retrieving tuples:
 - File scan, or index that matches a selection (in the query)
- ❖ A tree index matches (a conjunction of) terms that involve only attributes in a *prefix* of the search key.
 - E.g., Tree index on $\langle a, b, c \rangle$ matches the selection $a=5$ AND $b=3$, and $a=5$ AND $b>6$, but not $b=3$.
- ❖ A hash index matches (a conjunction of) terms that has a term *attribute = value* for every attribute in the search key of the index.
 - E.g., Hash index on $\langle a, b, c \rangle$ matches $a=5$ AND $b=3$ AND $c=5$; but it does not match $b=3$, or $a=5$ AND $b=3$, or $a>5$ AND $b=3$ AND $c=5$.



One Approach to Selections

15

- Find the *most selective access path*, retrieve tuples using it, and apply any remaining terms that don't match the index:
 - *Most selective access path*: An index or file scan that we estimate will require the fewest page I/Os.
 - Terms that match this index reduce the number of tuples *retrieved*; other terms are used to discard some retrieved tuples, but do not affect number of tuples/pages fetched.
 - Consider $day < 8/9/94$ AND $bid=5$ AND $sid=3$. A B+ tree index on *day* can be used; then, $bid=5$ and $sid=3$ must be checked for each retrieved tuple. Similarly, a hash index on $\langle bid, sid \rangle$ could be used; $day < 8/9/94$ must then be ~~checked~~.



Using an Index for Selections

16

- Cost depends on #qualifying tuples, and clustering.
 - Cost of finding qualifying data entries (typically small) plus cost of retrieving records (could be large w/o clustering).
 - In example, assuming uniform distribution of names, about 10% of tuples qualify (100 pages, 10000 tuples). With a clustered index, cost is little more than 100 I/Os; if unclustered, upto 10000 I/Os!

```
SELECT *
FROM Reserves R
WHERE R.rname < 'C%'
```

11/16/11



Projection

```
SELECT DISTINCT
  R.sid, R.bid
FROM   Reserves R
```

17

- The expensive part is removing duplicates.
 - SQL systems don't remove duplicates unless the keyword **DISTINCT** is specified in a query.
- Sorting Approach: Sort on <sid, bid> and remove duplicates. (Can optimize this by dropping unwanted information while sorting.)
- Hashing Approach: Hash on <sid, bid> to create partitions. Load partitions into memory one at a time, build in-memory hash structure, and eliminate duplicates.
- If there is an index with both R.sid and R.bid in the search key, may be cheaper to use data entries!

11/16/11



Example

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating=8
```

18

- If we have an index on *rating*:
 - $(1/NKeys(I)) * NTuples(R) = (1/10) * 40000$ tuples retrieved.
 - Clustered index: $(1/NKeys(I)) * (NPages(I) + NPages(R)) = (1/10) * (50 + 500)$ pages are retrieved. (This is the **cost**.)
 - Unclustered index: $(1/NKeys(I)) * (NPages(I) + NTuples(R)) = (1/10) * (50 + 40000)$ pages are retrieved.
- If we have an index on *sid*:
 - Would have to retrieve all tuples/pages in file (500). Compare with a clustered index *rating*, the cost is 55, with unclustered index, 5+4000.
- Doing a file scan:
 - We retrieve all file pages (500).


11/16/11



Queries Over Multiple Relations

19

- Fundamental decision in System R: only left-deep join trees are considered.
 - As the number of joins increases, the number of alternative plans grows rapidly; we need to restrict the search space.
 - Left-deep trees allow us to generate all *fully pipelined* plans.
 - Intermediate results not written to temporary files.
 - Not all left-deep trees are fully pipelined (e.g., SM join).



11/16/11



Enumeration of Left-Deep Plans

20

- Left-deep plans differ only in: order of relations, access method for each relation, and join method for each join.
- Enumerated using N passes (if N relations joined):
 - Pass 1: Find best 1-relation plan for each relation.
 - Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (All 2-relation plans.)
 - Pass N: Find best way to join result of a (N-1)-relation plan (as outer) to the N'th relation. (All N-relation plans.)
- For each subset of relations, retain only:
 - Cheapest plan overall, plus
 - Cheapest plan for each interesting order of the tuples.

11/16/11



Enumeration of Plans (Contd.)

21

- ORDER BY, GROUP BY, aggregates etc. handled as a final step, using either an 'interestingly ordered' plan or an additional sorting operator.
- An N-I way plan is not combined with an additional relation unless there is a join condition between them, (the one exception being where all predicates in WHERE have been used up meaning the query contains a cartesian product)
 - i.e., avoid Cartesian products if possible.
- In spite of pruning plan space, this approach is still exponential in the # of tables.

11/16/11



Cost Estimation for Multi-relation Plans

22

```
SELECT attribute list
FROM relation list
WHERE term1 AND ... AND termk
```

- Consider a query block:
- Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause
- *Reduction factor (RF)* associated with each *term* reflects the impact of the *term* in reducing result size.
 - *Result cardinality* = Max # tuples * product of all RF's.
- Multi-relation plans are built up by joining one new relation at a time.
 - Cost of join method, plus estimation of join cardinality gives us both cost estimate and result size estimate.

11/16/11



```
SELECT sname
FROM Reserves R, Sailors S
WHERE R.sid = S.sid and
      bid = 100 and rating > 5
```

23

Pass 1:

- *Sailors*: B+ tree matches *rating > 5*, probably cheapest. However, if selection is expected to retrieve a lot of tuples, and index is unclustered, file scan may be cheaper.
 - *sid* is an interesting order, so hash on *sid* kept even if higher cost than *rating* index
- *Reserves*: B+ tree on *bid* matches *bid = 100*; cheapest.

Pass 2:

Consider each plan retained from Pass 1 as the outer, and consider how to join it with the (only) other relation.

Reserves as outer: Hash index can be used to get Sailors tuples that satisfy *sid* = outer tuple's *sid* value (selection on *rating* moved **after** join)
Alternative is BNL with $\sigma_{rating > 5}(Sailors)$

Sailors as outer: block-nested loop to join with $\sigma_{bid=100}(Reserves)$

11/16/11




Sailors:

B+ tree on *rating*
Hash on *sid*

Reserves:

B+ tree on *bid*



Nested Queries

24

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS
  (SELECT *
   FROM Reserves R
   WHERE R.bid=103
        AND R.sid=S.sid)
```

- Nested block is optimized independently, with the outer tuple considered as providing a selection condition.
- Outer block is optimized with the cost of 'calling' nested block computation taken into account.
- Implicit ordering of these blocks means that some good strategies are not considered. *The non-nested version of the query is typically optimized better.*

Nested block to optimize:

```
SELECT *
FROM Reserves R
WHERE R.bid=103
      AND R.sid= outer value
```

Equivalent non-nested query:

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid
      AND R.bid=103
```

11/16/11



Summary

25

- Query optimization is an important task in a relational DBMS.
- Must understand optimization in order to understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).
- Two parts to optimizing a query:
 - ▣ Consider a set of alternative plans.
 - Must prune search space; typically, left-deep plans only.
 - ▣ Must estimate cost of each plan that is considered.
 - Must estimate size of result and cost for each plan node.
 - Key issues: Statistics, indexes, operator implementations.

11/16/11



Summary (Contd.)

26

- Single-relation queries:
 - ▣ All access paths considered, cheapest is chosen.
 - ▣ Issues: Selections that *match* index, whether index key has all needed fields and/or provides tuples in a desired order.
- Multiple-relation queries:
 - ▣ All single-relation plans are first enumerated.
 - Selections/projections considered as early as possible.
 - ▣ Next, for each 1-relation plan, all ways of joining another relation (as inner) are considered.
 - ▣ Next, for each 2-relation plan that is 'retained', all ways of joining another relation (as inner) are considered, etc.
 - ▣ At each level, for each subset of relations, only best plan for each interesting order of tuples is 'retained'.

