



CS 443 Recovery Chapter 18

Slides adapted from
Ramakrishnan & Gerhke
pages.cs.wisc.edu/~dbbook/

Adapted from Cow Book 3rd Ed.

Review: The ACID properties

2

- ❖ **A** tomicity: All actions in the Xact happen, or none happen.
- ❖ **C** onsistency: If each Xact is consistent, and the DB starts consistent, it ends up consistent.
- ❖ **I** solation: Execution of one Xact is isolated from that of other Xacts.
- ❖ **D** urability: If a Xact commits, its effects persist.

□ The **Recovery Manager** guarantees Atomicity & Durability.

11/16/11




Motivation

3

- Atomicity:
 - Transactions may abort (“Rollback”).
- Durability:
 - What if DBMS stops running? (Causes?)

❖ Desired Behavior after system restarts:

- T1, T2 & T3 should be durable.
- T4 & T5 should be aborted (effects not seen).



11/16/11



Assumptions

4

- Concurrency control is in effect.
 - Strict 2PL, in particular.
- Updates are happening “in place”.
 - i.e. data is overwritten on (deleted from) the disk.
- A simple scheme to guarantee Atomicity & Durability?

11/16/11



Handling the Buffer Pool

5

- Force every write to disk on commit?

- ▣ Poor response time.
 - ▣ But provides durability.

- Steal buffer-pool frames from uncommitted Xacts?

- ▣ If not, poor throughput.
 - ▣ If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial	
No Force		Desired

11/16/11



More on Steal and Force

6

- **STEAL** (why enforcing Atomicity is hard)

- ▣ To steal frame *F*: Current page in *F* (say *P*) is written to disk; some Xact holds lock on *P*.

- ⊠ What if the Xact with the lock on *P* aborts?

- ⊠ Must remember the old value of *P* at steal time (to support UNDOing the write to page *P*).

- **NO FORCE** (why enforcing Durability is hard)

- ▣ What if system crashes before a modified page is written to disk?

- ▣ Write as little as possible, in a convenient place, at commit time, to support REDOing modifications.

11/16/11



Basic Idea: Logging

7



- Record REDO and UNDO information, for every update, in a *log*.

- ▣ Sequential writes to log (put it on a separate disk).
 - ▣ Minimal info (diff) written to log, so multiple updates fit in a single log page.

- **Log**: An ordered list of REDO/UNDO actions

- ▣ Log record contains:

<XID, pageID, offset, length, old data, new data>

- ▣ and additional control info (which we'll see soon).

11/16/11



Write-Ahead Logging (WAL)

8

- The Write-Ahead Logging Protocol:

- Must force the log record for an update before the corresponding data page gets to disk.
 - Must write all log records for a Xact before commit.

- #1 guarantees Atomicity.

- #2 guarantees Durability.

- Exactly how is logging (and recovery!) done?

- ▣ We'll study the ARIES algorithms.

11/16/11



WAL & the Log



- Each log record has a unique Log Sequence Number (LSN).

- LSNs always increasing.

- Each data page contains a pageLSN.


- The LSN of the most recent *log record* for an update to that page.

- System keeps track of flushedLSN.

- The max LSN flushed so far.

- WAL: Before a page is written,


- $\text{pageLSN} \leq \text{flushedLSN}$



Log Records

10

LogRecord fields:



Possible log record types:

- Update**
- Commit**
- Abort**
- End** (signifies end of commit or abort)
- Compensation Log Records (CLRs)
 - for UNDO actions

11/16/11



Other Log-Related State

11

- Transaction Table:

- One entry per active Xact.

- Contains XID, status (running/committed/aborted), and lastLSN.

- Dirty Page Table:

- One entry per dirty page in buffer pool.

- Contains recLSN -- the LSN of the log record which **first** caused the page to be dirty.

11/16/11



Normal Execution of an Xact

12

- Series of reads & writes, followed by commit or abort.

- We will assume that write (of a page) is atomic on disk.
 - In practice, additional details to deal with non-atomic writes.

- Strict 2PL.

- STEAL, NO-FORCE buffer management, with Write-Ahead Logging.

11/16/11



Checkpointing


13

- Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. Write to log:
 - **begin_checkpoint** record: Indicates when chkpt began.
 - **end_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a 'fuzzy checkpoint':
 - ☒ Other Xacts continue to run; so these tables accurate only as of the time of the begin_checkpoint record.
 - ☒ No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page. (So it's a good idea to periodically flush dirty pages to disk!)
 - Store LSN of chkpt record in a safe place (**master record**).




The Big Picture: What's Stored Where

14



LogRecords

prevLSN
XID
type
pageID
length
offset
before-image
after-image



Data pages
each
with a
pageLSN

master record



Xact Table

XID
lastLSN
status

Dirty Page Table

recLSN

flushedLSN

11/16/11



Simple Transaction Abort

15


- For now, consider an explicit abort of a Xact.
 - No crash involved.
- We want to "play back" the log in reverse order, UNDOing updates.
 - Get lastLSN of Xact from Xact table.
 - Can follow chain of log records backward via the prevLSN field.
 - Before starting UNDO, write an *Abort* log record.
 - ☒ For recovering from crash during UNDO!

11/16/11



Abort, cont.

16



- To perform UNDO, must have a lock on data!
 - No problem!
- Before restoring old value of a page, write a CLR:
 - You continue logging while you UNDO!!
 - CLR has one extra field: undonextLSN
 - ☒ Points to the next LSN to undo (i.e. the prevLSN of the record we're currently undoing).
 - CLR's never Undone (but they might be Redone when repeating history: guarantees Atomicity!)
- At end of UNDO, write an "end" log record.

11/16/11



Transaction Commit

17


- Write commit record to log.
- All log records up to Xact's lastLSN are flushed.
 - Guarantees that flushedLSN \geq lastLSN.
 - Note that log flushes are sequential, synchronous writes to disk.
 - Many log records per log page.
- Commit() returns.
- Write end record to log.

11/16/11



Crash Recovery: Big Picture

18



- ❖ Start from a checkpoint (found via master record).
- ❖ Three phases. Need to:
 - Figure out which Xacts committed since checkpoint, which failed (Analysis).
 - REDO *all* actions.
 - ◆ (repeat history)
 - UNDO effects of failed Xacts.

11/16/11



Recovery: The Analysis Phase

19

- Reconstruct state at checkpoint.
 - via end_checkpoint record.
- Scan log forward from checkpoint.
 - End record: Remove Xact from Xact table.
 - Other records: Add Xact to Xact table, set lastLSN=LSN, change Xact status on commit.
 - Update record: If P not in Dirty Page Table,
 - ☒ Add P to D.P.T., set its recLSN=LSN.

11/16/11



Recovery: The REDO Phase

20

- We repeat *History* to reconstruct state at crash:
 - Reapply *all* updates (even of aborted Xacts!), redo CLR's.
- Scan forward from log rec containing smallest recLSN in D.P.T. For each CLR or update log rec LSN, REDO the action unless:
 - Affected page is not in the Dirty Page Table, or
 - Affected page is in D.P.T., but has recLSN $>$ LSN, or
 - pageLSN (in DB) \geq LSN.
- To REDO an action:
 - Reapply logged action.
 - Set pageLSN to LSN. No additional logging!

11/16/11



Recovery: The UNDO Phase

21

ToUndo = { / / / a lastLSN of a "loser" Xact }

Repeat:

- ❑ Choose largest LSN among ToUndo.
- ❑ If this LSN is a CLR and undonextLSN == NULL
 - ⊠ Write an End record for this Xact.
- ❑ If this LSN is a CLR, and undonextLSN != NULL
 - ⊠ Add undonextLSN to ToUndo
- ❑ Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo. if undonextLSN == NULL write End log record

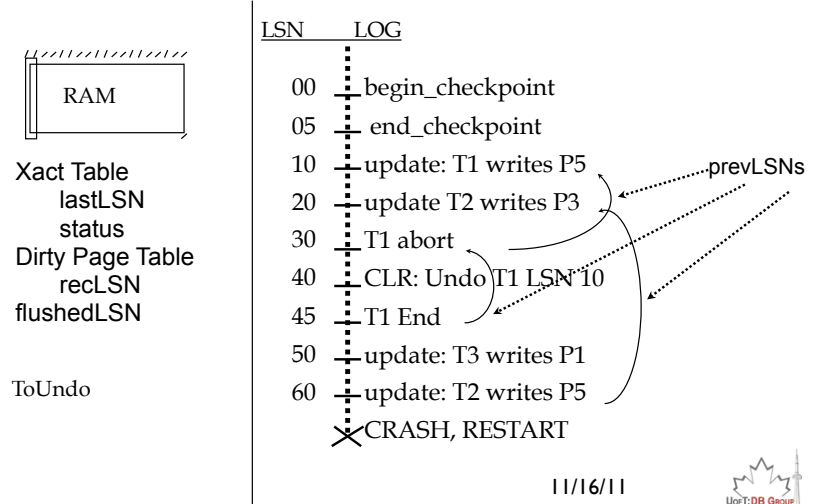
Until ToUndo is empty.

11/16/11



Example of Recovery

22

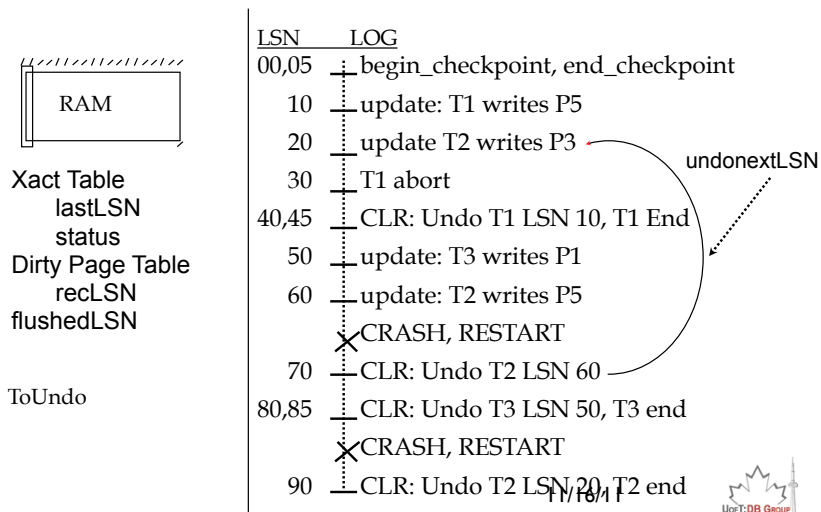


11/16/11



Example: Crash During Restart!

23



11/16/11



Additional Crash Issues

24

- ❑ What happens if system crashes during Analysis? During REDO?
- ❑ How do you limit the amount of work in REDO?
 - ❑ Flush asynchronously in the background.
 - ❑ Watch "hot spots"!
- ❑ How do you limit the amount of work in UNDO?
 - ❑ Avoid long-running Xacts.

11/16/11



Summary of Logging/Recovery

25

- Recovery Manager guarantees Atomicity & Durability.
- Use WAL to allow STEAL/NO-FORCE w/o sacrificing correctness.
- LSNs identify log records; linked into backwards chains per transaction (via prevLSN).
- pageLSN allows comparison of data page and log records.

11/16/11



Summary, Cont.

26

- Checkpointing: A quick way to limit the amount of log to scan on recovery.
- Recovery works in 3 phases:
 - Analysis: Forward from checkpoint.
 - Redo: Forward from oldest recLSN.
 - Undo: Backward from end to first LSN of oldest Xact alive at crash.
- Upon Undo, write CLR.
- Redo “repeats history”: Simplifies the logic!

11/16/11

